

# Lecture 7: Generalization Theory, Model Assessment & Selection

Readings: Bach (Ch. 4), ESL (Ch. 7), ISL (Ch. 5); code

Soon Hoe Lim

September 16, 2025

# Outline

- ① Part I: Theoretical Foundations
- ② Part II: Practical Model Assessment & Selection
- ③ Other Measures of Complexity
- ④ Summary
- ⑤ Exercises

# Overview

- ❗ How to understand and build models that generalize well to unseen data?

## Part I: Theoretical Foundations (Bach Ch. 4)

**Goal:** Understand *why* certain models generalize better through statistical learning theory.

- ▶ Convex surrogates for intractable problems
- ▶ Risk decomposition and the sources of error
- ▶ Rademacher complexity for generalization bounds

## Part II: Practical Methods (ESL Ch. 7, ISL Ch. 5)

**Goal:** Learn *how* to assess and select models in practice.

- ▶ Bias-variance tradeoff and model complexity
- ▶ Cross-validation and resampling methods
- ▶ Information criteria (AIC, BIC) for model selection

# The Learning Problem: Formal Setup

## Definition 1: Statistical Learning Framework

- ▶ **Data:**  $(x_1, y_1), \dots, (x_N, y_N)$  drawn i.i.d. from unknown distribution  $P$  on  $\mathcal{X} \times \mathcal{Y}$
- ▶ **Goal:** Learn function  $f : \mathcal{X} \rightarrow \mathcal{Y}$  that minimizes expected risk:

$$R(f) = \mathbb{E}_{(x,y) \sim P}[\ell(y, f(x))]$$

- ▶ **Bayes Optimal:**  $f^* = \arg \min_f R(f)$  with  $R^* = R(f^*)$
- ▶ **Our Focus:** Methods based on **Empirical Risk Minimization (ERM)**:

$$\hat{f} = \arg \min_{f \in \mathcal{F}} \hat{R}(f) = \arg \min_{f \in \mathcal{F}} \frac{1}{N} \sum_{i=1}^N \ell(y_i, f(x_i))$$

**Key Challenge:** For binary classification with 0-1 loss, this is a combinatorial optimization problem. We need better approaches!

# The Convexification Strategy

For binary classification ( $\mathcal{Y} = \{-1, 1\}$ ) with 0-1 loss:

1. Learn real-valued function  $g : \mathcal{X} \rightarrow \mathbb{R}$  using convex surrogate loss  $\Phi(yg(x))$
2. Make predictions via  $f(x) = \text{sign}(g(x))$

## Benefits:

- ▶ **Computational:** Convex optimization (global optimum, efficient algorithms)
- ▶ **Theoretical:** Enables Rademacher complexity analysis
- ▶ **Practical:** Allows gradient-based methods and principled regularization


For binary classification, consider surrogate losses of the form  $\Phi(yg(x))$  where  $yg(x)$  is the **margin**.

- ▶ **Margin**  $> 0$ : Correct classification with confidence
- ▶ **Margin**  $< 0$ : Misclassification

# Convexification with Calibrated Surrogate Losses

A surrogate loss  $\Phi$  is **calibrated** if minimizing  $\mathbb{E}[\Phi(yg(x))]$  leads to the Bayes optimal classifier. All losses below are calibrated:

- ▶ **Hinge Loss (SVM):**  $\Phi(u) = \max(0, 1 - u)$
- ▶ **Logistic Loss:**  $\Phi(u) = \log(1 + e^{-u})$
- ▶ **Exponential Loss (AdaBoost):**  $\Phi(u) = e^{-u}$
- ▶ **Squared Loss:**  $\Phi(u) = (1 - u)^2$



# Risk Decomposition: Sources of Error

We consider loss functions that are defined for real-valued outputs (for binary classification problems we will use a surrogate loss).

For any ERM estimator  $\hat{f} \in \mathcal{F}$  trained on  $N$  samples:

$$\mathbb{E}[R(\hat{f})] - R^* = \underbrace{\mathbb{E}[R(\hat{f})] - \inf_{f \in \mathcal{F}} R(f)}_{\text{Estimation Error}} + \underbrace{\inf_{f \in \mathcal{F}} R(f) - R^*}_{\text{Approximation Error}}$$

- ▶ **Approximation Error:** Fundamental limitation of the model class  $\mathcal{F}$ 
  - ▶ How well can the *best possible* function in  $\mathcal{F}$  approximate  $f^*$ ?
  - ▶ Only reduced by choosing more flexible  $\mathcal{F}$  (more complex models)
  - ▶ Independent of sample size  $N$
- ▶ **Estimation Error:** Error from finite sample learning
  - ▶ How much worse is our empirical solution  $\hat{f}$  vs. the best in class?
  - ▶ Decreases with more data (typically  $O(1/\sqrt{N})$  or better)
  - ▶ Increases with model complexity (richer  $\mathcal{F}$ )

# Uniform Deviation

Estimation error is controlled by uniform deviation over the function class:

$$\mathbb{E}[R(\hat{f})] - \inf_{f \in \mathcal{F}} R(f) \leq 2\mathbb{E} \left[ \sup_{f \in \mathcal{F}} |R(f) - \hat{R}(f)| \right]$$

## Proposition 1: Application of McDiarmid's Inequality

Let  $Z = \sup_{f \in \mathcal{F}} (R(f) - \hat{R}(f))$  where loss is bounded:  $\ell(y, f(x)) \in [0, \ell_\infty]$ . Then,

$$P(Z \geq \mathbb{E}[Z] + t) \leq \exp \left( -\frac{2Nt^2}{\ell_\infty^2} \right),$$

**Proof.**

See blackboard. □

**The Problem:** We still need to bound  $\mathbb{E}[\sup_{f \in \mathcal{F}} |R(f) - \hat{R}(f)|]$ . This is where Rademacher complexity comes in.



# Rademacher Complexity: Definition

## Definition 2: Rademacher Complexity

Consider  $N$  i.i.d. sample  $z_1, \dots, z_N \in \mathcal{Z}$  and a class  $\mathcal{H}$  of functions from  $\mathcal{Z}$  to  $\mathbb{R}$ . The Rademacher complexity of  $\mathcal{H}$  is:

$$\mathcal{R}_N(\mathcal{H}) = \mathbb{E}_{\mathbf{z}, \varepsilon} \left[ \sup_{h \in \mathcal{H}} \frac{1}{N} \sum_{i=1}^N \varepsilon_i h(z_i) \right]$$

where  $\varepsilon_i \stackrel{\text{iid}}{\sim} \text{Rademacher}(\pm 1)$  are independent of the data.

- ▶ **Random Noise Test:** How well can functions in  $\mathcal{H}$  fit random noise ( $\varepsilon_i$ )?
- ▶ **Complexity Measure:** Rich function classes can fit noise well  $\Rightarrow$  high  $\mathcal{R}_N(\mathcal{H})$
- ▶ **Sample Size Effect:** More data makes it harder to fit noise  $\Rightarrow \mathcal{R}_N(\mathcal{H})$  typically decreases as  $N$  increases
- ▶ **Dimension Independent:** Often gives dimension-free bounds

# Symmetrization: The Key Lemma

**Connection to Generalization:** Functions that can fit random noise well are prone to overfitting real data. Rademacher complexity quantifies this overfitting potential.

A useful lemma:

## Lemma 1: Symmetrization Lemma

For any function class  $\mathcal{H}$ :

$$\mathbb{E} \left[ \sup_{h \in \mathcal{H}} \left( \frac{1}{N} \sum_{i=1}^N h(z_i) - \mathbb{E}[h(Z)] \right) \right] \leq 2\mathcal{R}_N(\mathcal{H}).$$

Proof.

See blackboard. □

# Rademacher Generalization Bound

## Theorem 1: Rademacher Generalization Bound

Let  $\mathcal{F}$  be a function class and  $\ell$  be a loss function with  $|\ell(y, f(x))| \leq M$ . Define the loss function class  $\mathcal{H} = \{(x, y) \mapsto \ell(y, f(x)) : f \in \mathcal{F}\}$ . Then for any  $\delta > 0$ , with probability at least  $1 - \delta$ ,

$$\sup_{f \in \mathcal{F}} (R(f) - \hat{R}(f)) \leq 2\mathcal{R}_N(\mathcal{H}) + M \sqrt{\frac{\log(1/\delta)}{2N}}.$$

Proof.

See blackboard. □

## Proposition 2: Contraction Principle

If  $\ell(y, \cdot)$  is  $L$ -Lipschitz for each  $y$ , then for loss class  $\mathcal{H} = \{\ell(\cdot, f(\cdot)) : f \in \mathcal{F}\}$ ,  $\mathcal{R}_N(\mathcal{H}) \leq L \cdot \mathcal{R}_N(\mathcal{F})$ .

Proof.

See blackboard. □

# Rademacher Complexity: Key Properties

## Proposition 3: Linear Functions

For  $\mathcal{F} = \{f_\theta(x) = \theta^\top \phi(x) : \|\theta\|_2 \leq D\}$ :

$$\mathcal{R}_N(\mathcal{F}) = \frac{D}{N} \mathbb{E} \left[ \left\| \sum_{i=1}^N \varepsilon_i \phi(x_i) \right\|_2 \right]$$

If  $\mathbb{E}[\|\phi(X)\|_2^2] \leq R^2$ , then:  $\mathcal{R}_N(\mathcal{F}) \leq \frac{DR}{\sqrt{N}}$ .

## Proof.

See blackboard. □

💡 **Dimension-Free Bound:** The bound  $DR/\sqrt{N}$  does not depend on the dimension of  $\phi(x)$ ! Applicable to infinite-dimensional spaces (e.g., RKHS).

# Main Generalization Theorem

## Theorem 2: Generalization Bound for Linear Predictors

Consider:

- ▶ Linear predictors:  $\mathcal{F} = \{f_\theta(x) = \theta^\top \phi(x) : \|\theta\|_2 \leq D\}$
- ▶ Loss function is  $L$ -Lipschitz and bounded
- ▶ Bounded features:  $\mathbb{E}[\|\phi(X)\|_2^2] \leq R^2$
- ▶ ERM estimator  $\hat{f} = \arg \min_{f \in \mathcal{F}} \hat{R}(f)$

Then  $\mathbb{E}[R(\hat{f})] - \inf_{f \in \mathcal{F}} R(f) \leq \frac{4LDR}{\sqrt{N}}$ . Moreover, for any  $\delta > 0$ , with probability  $\geq 1 - \delta$ ,  $R(\hat{f}) - \inf_{f \in \mathcal{F}} R(f) \leq \frac{4LDR}{\sqrt{N}} + \ell_\infty \sqrt{\frac{\log(1/\delta)}{2N}}$ .

Proof.

See blackboard. □

- ▶ Rate  $O(1/\sqrt{N})$  is **minimax optimal** for this setting
- ▶ Bound scales with  $L \cdot D \cdot R$  (loss smoothness  $\times$  model complexity  $\times$  data scale)
- ▶ **Dimension-free:** Works for infinite-dimensional  $\phi(x)$  (kernels!)
- ▶ Others: VC dimension bound, PAC-Bayes bound (not covered)


# From Theory to Practice

Linear models can achieve good theoretical guarantees. But in practice:

- ▶ We have **multiple candidate models** (LDA, logistic regression, SVM, etc.)
- ▶ Models have **tuning parameters** (regularization strength  $\lambda$ , kernel bandwidth, etc.)
- ▶ We need **reliable estimates** of how well our final model will perform

## Definition 3: Model Assessment vs Model Selection

- ▶ **Model Selection:** Choose the best model from a set of candidates by comparing their estimated test performance
- ▶ **Model Assessment:** Estimate the test error of our final chosen model to understand its real-world performance

 **Golden Rule:** The same data should not be used for both model selection and final assessment! This leads to overly optimistic performance estimates.

# The Fundamental Goal: Minimizing Test Error

## Definition 4: Training vs Test Error

Given training set  $\mathcal{T} = \{(x_1, y_1), \dots, (x_N, y_N)\}$  and learned model  $\hat{f}$ :

- **Training Error:**  $\text{err}_{\mathcal{T}} = \frac{1}{N} \sum_{i=1}^N L(y_i, \hat{f}(x_i))$

- **Test Error (Generalization Error):**  $\text{Err}_{\mathcal{T}} = \mathbb{E}_{X^0, Y^0} [L(Y^0, \hat{f}(X^0)) \mid \mathcal{T}]$   
where  $(X^0, Y^0)$  is a new test point independent of  $\mathcal{T}$ .

**The Core Challenge:** We want to minimize test error, but we can only observe training error. Low training error low test error due to **overfitting**.

► **Regression:**  $L(y, \hat{f}(x)) = (y - \hat{f}(x))^2$  (squared error)

► **Classification:**  $L(y, \hat{f}(x)) = \mathbb{I}[y \neq \hat{f}(x)]$  (0-1 loss)

# The Optimism of Training Error

## Definition 5: Optimism

The **optimism** is the expected difference between test and training error:  $\text{op} = \mathbb{E}_{\mathcal{T}}[\text{Err}_{\mathcal{T}}] - \mathbb{E}_{\mathcal{T}}[\text{err}_{\mathcal{T}}]$ .

This gives us the fundamental relationship:  $\mathbb{E}_{\mathcal{T}}[\text{Err}_{\mathcal{T}}] = \mathbb{E}_{\mathcal{T}}[\text{err}_{\mathcal{T}}] + \text{op}$

**Optimism for Linear Models.** For a linear model with  $d$  parameters fitted by least squares:  $\text{op} = \frac{2d}{N} \sigma^2$ .

💡 Each parameter "uses up"  $2\sigma^2/N$  worth of degrees of freedom.

More complex models (larger  $d$ ) have higher optimism:

- ▶ **Linear model:**  $\text{op} \propto d$
- ▶ **Neural networks:**  $\text{op} \propto \# \text{ parameters}$
- ▶ **Nonparametric methods:**  $\text{op} \propto \text{effective degrees of freedom}$



# Information Criteria: AIC, BIC, and $C_p$

These methods estimate test error by correcting the training error with a complexity penalty.

## Definition 6: General Form

$$\widehat{\text{Err}} = \text{Training Error} + \text{Penalty}(\text{Model Complexity})$$

## Specific Criteria

- ▶ **Mallows'  $C_p$  (for OLS regression):**  $C_p = \text{err}_{\mathcal{T}} + \frac{2d}{N}\hat{\sigma}^2$  where  $\hat{\sigma}^2$  is noise variance estimate (typically from full model).
- ▶ **Akaike Information Criterion (AIC):**  $\text{AIC} = -2 \log \mathcal{L}(\hat{\theta}) + 2d$  For Gaussian errors,  $-2 \log \mathcal{L} \propto \text{RSS}$ , so  $\text{AIC} \approx C_p$ .
- ▶ **Bayesian Information Criterion (BIC):**  $\text{BIC} = -2 \log \mathcal{L}(\hat{\theta}) + d \log N$  Since  $\log N \geq 2$  for  $N \geq 8$ , BIC penalizes complexity more than AIC.

# Information Criteria: Properties and Usage

## Theoretical Properties

- ▶ **AIC:** Asymptotically equivalent to cross-validation
  - ▶ Tends to select models with complexity close to optimal for prediction
  - ▶ Can overfit in small samples
- ▶ **BIC:** Consistent model selection
  - ▶ If true model is in candidate set, BIC selects it with probability  $\rightarrow 1$
  - ▶ More conservative (selects simpler models) than AIC
  - ▶ Better for interpretation, worse for prediction when truth is complex

## Practical Guidelines

- ▶ **Use AIC** when primary goal is prediction accuracy
- ▶ **Use BIC** when seeking to identify "true" parsimonious model
- ▶ Both require likelihood-based models and proper parameter counting
- ▶ Not directly applicable to non-likelihood methods (SVM, trees, etc.)

⚠ **Limitation:** Information criteria rely on asymptotic approximations and may be unreliable with small samples or misspecified models.

# Cross-Validation: The Gold Standard

Cross-validation directly estimates test error by simulating the train/test process:

- ▶ Split data into training and validation sets
- ▶ Train on training set, evaluate on validation set
- ▶ Repeat with different splits to get stable estimate

## K-Fold Cross-Validation

- 1: **Input:** Dataset  $\mathcal{D} = \{(x_1, y_1), \dots, (x_n, y_n)\}$ , number of folds  $K$
- 2: Randomly partition  $\mathcal{D}$  into  $K$  disjoint folds:  $\mathcal{D} = \mathcal{D}_1 \cup \dots \cup \mathcal{D}_K$
- 3: **for**  $k = 1, 2, \dots, K$  **do**
- 4:    $\mathcal{D}_{\text{train}}^{(k)} \leftarrow \mathcal{D} \setminus \mathcal{D}_k$  ▷ Training = all except fold  $k$
- 5:    $\mathcal{D}_{\text{val}}^{(k)} \leftarrow \mathcal{D}_k$  ▷ Validation = fold  $k$
- 6:   Train model:  $\hat{f}^{(k)} \leftarrow \text{Learn}(\mathcal{D}_{\text{train}}^{(k)})$
- 7:   Compute validation error:  $\text{Err}_k \leftarrow \frac{1}{|\mathcal{D}_k|} \sum_{(x,y) \in \mathcal{D}_k} L(y, \hat{f}^{(k)}(x))$
- 8: **end for**
- 9: **Return:**  $\text{CV}_K \leftarrow \frac{1}{K} \sum_{k=1}^K \text{Err}_k$

# Cross-Validation: Variants and Properties

- ▶ **5-fold or 10-fold CV:** Good bias-variance tradeoff, computationally feasible. 10-fold CV most common.
- ▶ **Leave-One-Out CV (LOOCV):**  $K = N$ , nearly unbiased but high variance
- ▶ **Stratified CV:** Maintains class proportions in each fold (for classification)
- ▶ **Time series CV:** Respects temporal order (no future information leak)

## Bias-Variance of CV Estimates

### Bias:

- ▶ CV uses  $(N - N/K)$  training samples vs.  $N$  for final model
- ▶ Small  $K$  (few folds)  $\rightarrow$  more bias (pessimistic estimates)
- ▶ Large  $K$  (many folds)  $\rightarrow$  less bias

### Variance:

- ▶ Small  $K \rightarrow$  less variance (fewer, more different estimates to average)
- ▶ Large  $K \rightarrow$  more variance (many, highly correlated estimates)
- ▶ LOOCV often has very high variance

# Cross-Validation for Model Selection

## Model Selection via Cross-Validation

- 1: **Input:** Dataset  $\mathcal{D}$ , candidate models  $\mathcal{M}_1, \dots, \mathcal{M}_m$
- 2: **for** each model  $\mathcal{M}_j$  **do**
- 3:     Compute  $\text{CV}_K(\mathcal{M}_j)$  using  $K$ -fold cross-validation
- 4: **end for**
- 5:  $\hat{j} \leftarrow \arg \min_j \text{CV}_K(\mathcal{M}_j)$  ▷ Select best model
- 6: **Return:** Best model  $\mathcal{M}_{\hat{j}}$


When both selecting and assessing models:

1. **Outer loop:** Split data into train/test
2. **Inner loop:** Use CV on training set to select best model
3. **Assessment:** Evaluate selected model on held-out test set

This prevents **selection bias** - overly optimistic estimates from using the same data for both selection and assessment.

⚠ **Common Mistake:** Using the CV score from model selection as the final performance estimate. This is biased! Need separate test set or nested CV.

# Example: Selecting A Polynomial Regression Model



Degree	AIC	BIC	CV_MSE
0	1335.76	1342.36	46.30
1	891.15	901.04	4.97
2	794.39	807.58	3.15
3	580.27	596.76	1.06
4	566.52	<b>586.31</b>	1.00
5	565.40	588.49	1.00
6	567.21	593.60	1.00
7	568.14	597.82	1.00
8	<b>562.42</b>	595.40	<b>0.96</b>

# The Bootstrap

## Bootstrap Principle

**Idea:** If we can't get new samples from the population, create "new" samples by resampling from our data.

### Bootstrap Procedure

- 1: **Input:** Original dataset  $\mathcal{T} = \{(x_1, y_1), \dots, (x_N, y_N)\}$
- 2: **for**  $b = 1, 2, \dots, B$  **do**
- 3:     Create  $\mathcal{T}^{*b}$  by sampling  $N$  points from  $\mathcal{T}$  **with replacement**
- 4:     Train model:  $\hat{f}^{*b} \leftarrow \text{Learn}(\mathcal{T}^{*b})$
- 5:     Compute statistic of interest on  $\hat{f}^{*b}$
- 6: **end for**
- 7: Analyze distribution of statistics across bootstrap samples

❗ Naive bootstrap (biased) vs. leave-one-out bootstrap

# The Bootstrap

The bootstrap is another resampling method for estimating risk. It involves drawing  $B$  "bootstrap samples" of size  $N$  from the training data *with replacement*.

Each bootstrap sample omits, on average, 36.8% of the original data points. We can use these "out-of-bag" (OOB) points to form an error estimate.

The **leave-one-out bootstrap** error estimate is:

$$\hat{R}_{boot} = \frac{1}{N} \sum_{i=1}^N \frac{1}{|C^{-i}|} \sum_{b \in C^{-i}} \ell(y_i, \hat{f}^{*b}(x_i))$$

where  $C^{-i}$  is the set of bootstrap samples that do not contain observation  $i$ .



# The .632 and .632+ Estimators

The leave-one-out bootstrap estimate can be biased upwards because models are trained on smaller, less diverse datasets. The .632 estimators correct for this.

## Definition 7: The .632 Estimator

A weighted average of the (optimistic) empirical risk and the (pessimistic) LOO bootstrap risk:

$$\hat{R}^{(.632)} = 0.368 \cdot \hat{R}_N(\hat{f}) + 0.632 \cdot \hat{R}_{boot}$$

This can fail in heavily overfit situations where  $\hat{R}_N(\hat{f}) \approx 0$ .

## Definition 8: The .632+ Estimator

An improved version that adapts to the amount of overfitting.

$$\hat{R}^{(.632+)} = (1 - \hat{w}) \cdot \hat{R}_N(\hat{f}) + \hat{w} \cdot \hat{R}_{boot}$$

The weight  $\hat{w}$  adapts to the relative overfitting rate, putting more trust in the out-of-sample estimate when the model fits the training data too closely.

# Vapnik-Chervonenkis (VC) Dimension

For classification models, the VC dimension provides a theoretical measure of model complexity, or the "richness" of the hypothesis class  $\mathcal{F}$ .

## Definition 9: Shattering and VC Dimension

A class of functions  $\mathcal{F}$  **shatters** a set of  $N$  points if it can realize every possible binary labeling of those points. The **VC dimension**  $h$  of  $\mathcal{F}$  is the maximum number of points that can be shattered.

The VC dimension provides a probabilistic upper bound on the true risk:

$$R(\hat{f}) \leq \hat{R}_N(\hat{f}) + \sqrt{\frac{h(\log(2N/h) + 1) - \log(\eta/4)}{N}}$$

⚠ This bound is often too loose for practical model selection but is of great theoretical importance.

# Theoretical vs Practical: Bridging the Gap?

## Theoretical Insights:

- ▶ **Convex surrogates** make hard problems tractable
- ▶ **Rademacher complexity** gives dimension-free generalization bounds
- ▶ **Regularization** can improve rates from  $O(1/\sqrt{N})$  to  $O(1/N)$
- ▶ **Linear models** have strong theoretical guarantees

## Practical Tools:

- ▶ **Cross-validation** works broadly without strong assumptions
- ▶ **Information criteria** efficient for likelihood-based models
- ▶ **Bootstrap** provides flexible uncertainty quantification
- ▶ **Nested CV** prevents selection bias in assessment

# Method Comparison and Recommendations

Method	Pros	Cons
<b>AIC/BIC/<math>C_p</math></b>	Fast computation, well-established theory, good for linear models	Requires likelihood or OLS, assumes model is approximately correct, can be unreliable in small samples
<b>10-Fold CV</b>	Widely applicable, minimal assumptions, direct test error estimate	Computationally intensive ( $10\times$ cost), can be unstable with small datasets
<b>LOOCV</b>	Nearly unbiased for test error, deterministic result	Very high variance, extremely expensive, can be unstable
<b>Bootstrap</b>	More stable than LOOCV, provides uncertainty estimates, handles complex models well	More complex to implement, can still have bias issues, requires many bootstrap samples
<b>Validation Set</b>	Simple to understand and implement, fast	Reduces effective sample size, can be unreliable with small datasets

# Practical Guidelines

1. **For most applications:** Use 10-fold cross-validation
  - ▶ Excellent bias-variance tradeoff
  - ▶ Works with any learning algorithm
  - ▶ Widely accepted and understood
2. **For linear models with likelihood:** Consider AIC/BIC as faster alternatives
  - ▶ AIC for prediction-focused applications
  - ▶ BIC for model interpretation and parsimony
3. **For very small datasets ( $N < 100$ ):** Use LOOCV or bootstrap
  - ▶ Every sample counts
  - ▶ Higher variance acceptable given limited data
4. **For model selection + assessment:** Use nested CV or train/validation/test split
  - ▶ Prevents optimistic bias
  - ▶ Essential for honest performance reporting

## Exercise 7

1. **Calibration.** Show that the logistic loss and the exponential loss are calibrated. Construct an example of a non-calibrated surrogate loss.
2. **Rademacher Complexity.** Solve Exercise 4.8 in Bach.
3. **Massart's Lemma.** Solve Exercise 4.9 in Bach.
4. **Generalization Bound for Kernel Methods.** Extend Theorem 2 to kernel methods.
5. **Train Error is Smaller Than Test Error, On Average.** Solve Exercise 2.9 in ESL.
6. **Best Subset Analysis.** Solve Exercise 7.9 in ESL.